

University of Moratuwa
Department of Electronics and Telecommunication



EN3160 – Image Processing and Machine Vision
Report Activity 01
200356A

Q1.



Figure 1 Image after the intensity transform

```
import cv2 as cv
import matplotlib.pyplot as plt
import numpy as np

points = np.array([(1,1),(50,50),(51,100),(150,255),(151,150),(255,255)])
#print(len(points))
transform = np.linspace(0,0,1)

for i in range(0,int(len(points)/2)):
    t1 = np.linspace(points[0+2*i,1], points[1+2*i,1], points[1+2*i,0]-points[0+2*i,0]+1)
    transform = np.concatenate((transform , t1), axis=0).astype('uint8')
    #print(len(t1))
    #print(len(transform))

fig, ax = plt.subplots()
ax.plot(transform)
ax.set_xlabel('input intensity') #, $f(\mathbf{x})$
ax.set_ylabel('output intensity') #, $\mathbf{f}(\mathbf{x})$
ax.set_xlim(0,255)
ax.set_ylim(0,255)
ax.set_aspect('equal')
plt.savefig('transform1.png')
plt.show()

img = cv.imread('../emma.jpg', cv.IMREAD_GRAYSCALE)
cv.namedWindow('Emma Image', cv.WINDOW_NORMAL)
cv.imshow('Emma Image', img)
cv.waitKey(0)

transf = cv.LUT(img, transform)
cv.namedWindow('Emma Image', cv.WINDOW_NORMAL)
cv.imshow('Emma Image', transf)
cv.waitKey(0)
cv.destroyAllWindows()
```

Figure 2 Code for the given transformation

Q2.

a. White matter enhanced




Figure 3 Original image




Figure 4 White matter enhanced image




Figure 5 White matter enhancing transform

```
Question 2.a white matter

import cv2 as cv
import matplotlib.pyplot as plt
import numpy as np

points = np.array([(1,1),(100,50),(101,75),(255,255)])
#print(len(points))
transform = np.linspace(0,0,1)

for i in range(0,int(len(points)/2)):
    t1 = np.linspace(points[0+2*i,1], points[1+2*i,1], points[1+2*i,0]-points[0+2*i,0]+1)
    transform = np.concatenate((transform , t1), axis=0).astype('uint8')
    #print(len(t1))
    #print(len(transform))
```

Figure 6 White matter enhancing transform code

b. Gray matter enhanced



Figure 7 Original image



Figure 8 Gray matter enhanced image

```
points = np.array([(1,1),(75,25),(76,26),(200,225),(201,226),(255,255)])
#print(len(points))
transform = np.linspace(0,0,1)

for i in range(0,int(len(points)/2)):
    t1 = np.linspace(points[0+2*i,1], points[1+2*i,1], points[1+2*i,0]-points[0+2*i,0]+1)
    transform = np.concatenate((transform, t1), axis=0).astype('uint8')
    #print(len(t1))
#print(len(transform))
```

Figure 10 Gray matter enhancing transform code




Figure 9 Gray matter enhancing transform




Here, most black parts are not in gray matter.

Therefore, enhancement is done only for the region where gray matter is.

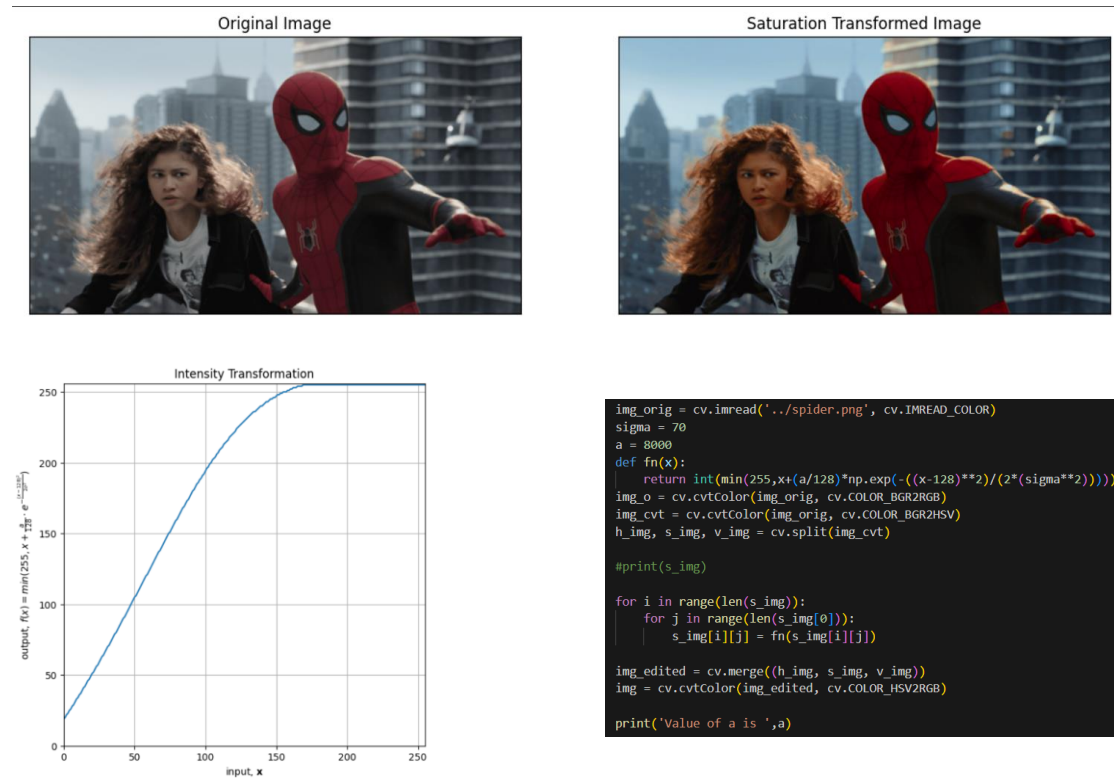
This intensity transform will enhance the gray matter area.

Q3.

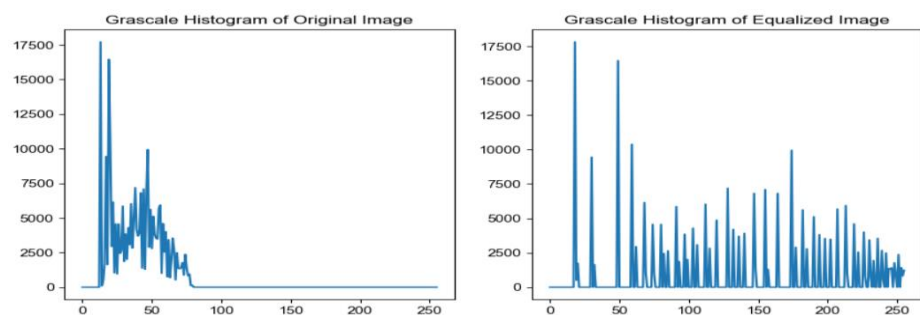
Gamma value used = 0.6



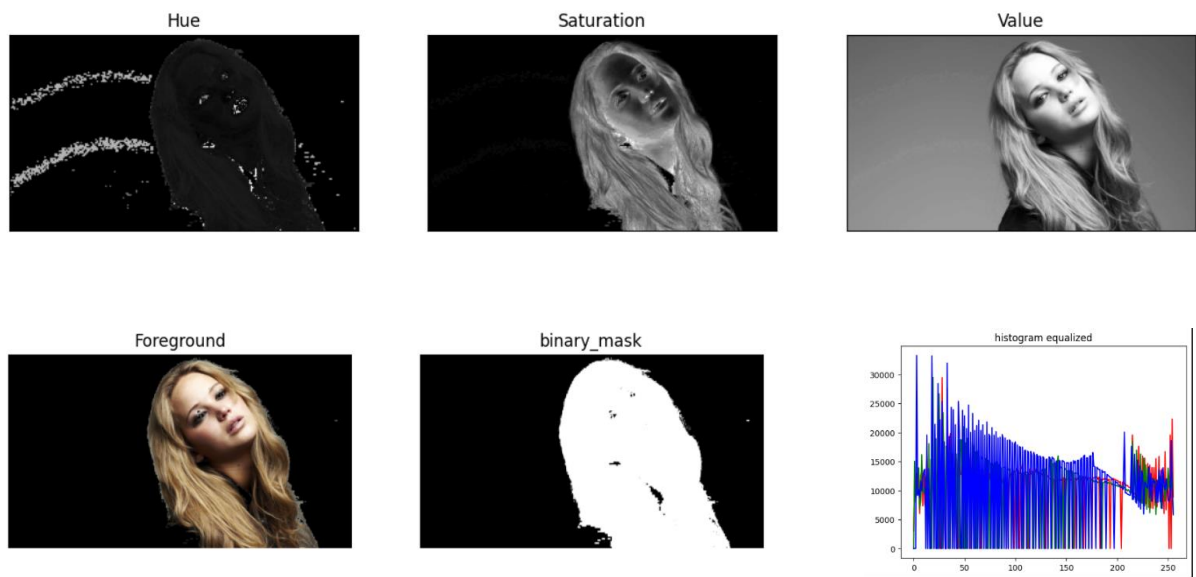
Q4.



Q5.



06.



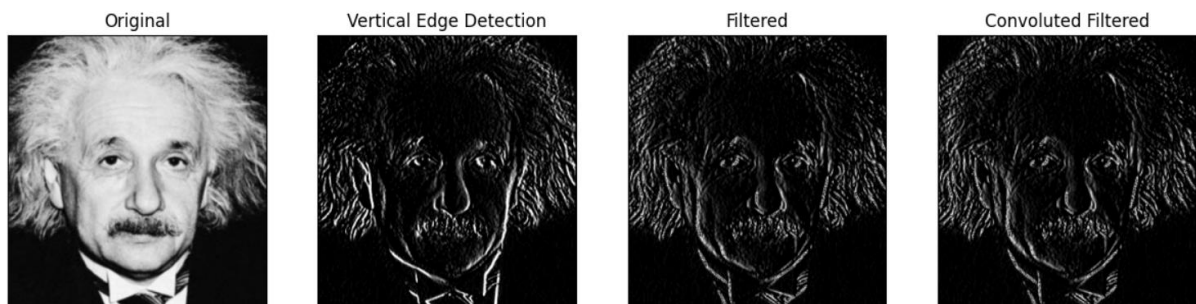
Histogram equalized foreground



Equalized foreground with normal background



Q7.



```
def k_filter(image, kernel):
    h,w = image.shape[:2]
    k_h,k_w = kernel.shape[0], kernel.shape[1]
    assert k_h%2 == 1 and k_w%2 == 1, "kernel is not accurate please check it"
    h_start, w_start = math.floor(k_h/2), math.floor(k_w/2)

    normalized_image = cv.normalize(image.astype('float'), None, 0.0, 1.0, cv.NORM_MINMAX)
    kernel_sum = np.sum(kernel)
    normalized_kernel = kernel

    #result image
    result = np.zeros(image.shape, dtype='float')

    #for convolution
    normalized_kernel = np.flipud(np.fliplr(normalized_kernel))

    for i in range(h_start, h-h_start):
        for j in range(w_start, w-w_start):
            result[i][j] = max(0,min(255.0,np.dot(normalized_image[i-h_start : i+h_start+1, j-w_start: j+w_start+1].flatten(), normalized_kernel.flatten()))
    #print(image)
    #print(result)
    result = 255.0 * result
    answer = result.astype(np.uint8)
    return answer
```



```
def conv(image, kernel):
    h,w = image.shape[:2]
    k,h,k_w = kernel.shape[0], kernel.shape[1]
    assert k_h%2 == 1 and k_w%2 == 1, "kernel is not accurate please check it"
    h_start, w_start = math.floor(k_h/2), math.floor(k_w/2)
    kernel = np.flipud(np.fliplr(kernel))

    #result image
    result = np.zeros(image.shape, dtype='float')

    for i in range(h_start, h-h_start):
        for j in range(w_start, w-w_start):
            result[i][j] = np.dot(image[i-h_start : i+h_start+1, j-w_start+1:j-w_start+1].flatten(), kernel.flatten())

    return result

img = cv.imread('.../einstein.png', cv.IMREAD_GRAYSCALE)

# Define a kernel
kernel = np.array([
    [1, 0, -1],
    [2, 0, -2],
    [1, 0, -1]
], dtype='float32')

# Apply 2D filtering with the defined kernel
imgc = cv.filter2D(img, -1, kernel)

fig, axes = plt.subplots(1, 4, sharex='all', sharey='all', figsize=(15, 15))
```

```
normalized_img = cv.normalize(img.astype('float'), None, 0.0, 1.0, cv.NORM_MINMAX)
conv1 = conv(normalized_img, kernel1)
#norm_img = np.maximum(0, np.minimum(255.0, conv2))
conv2 = conv(conv1, kernel2)
norm_img = np.maximum(0, np.minimum(255.0, conv2))

img_after = norm_img * 255
img_after = img_after.astype(np.uint8)
```

Q8.

```
def nearest_neighbour(image, scale):
    height, width, channels = image.shape
    h_zoom, w_zoom = height*scale, width*scale

    image_zoom = np.zeros((h_zoom, w_zoom, channels), dtype = np.float32)

    for h in range(h_zoom):
        for w in range(w_zoom):
            for c in range(channels):
                image_zoom[h][w][c] = image[int(math.floor(h/scale))][int(math.floor(w/scale))][c]

    return image_zoom.astype(np.uint8)
```

```
1 def bilinear_interpolation(image, scale):
2     height, width, channels = image.shape
3     h_zoom, w_zoom = height*scale, width*scale
4
5     image_zoom = np.zeros((h_zoom, w_zoom, channels), dtype = np.float32)
6
7     for h in range(h_zoom):
8         for w in range(w_zoom):
9             for c in range(channels):
10
11                 if ((h-1)%scale==0 and (w-1)%scale==0):
12                     image_zoom[h][w][c] = image[int((h-1)/scale) - 1][int((w-1)/scale) - 1][c]
13
14                 elif ((h-1)%scale==0):
15                     if (w%scale):
16                         image_zoom[h][w][c] = (w-1)/scale * image[int((h-1)/scale) - 1][int((w-1)/scale)][c]
17                     else:
18                         image_zoom[h][w][c] = (1 - (w-1)/scale) * image[int((h-1)/scale) - 1][int((w-1)/scale) - 1][c] + (w-1)/scale * image[int((h-1)/scale) - 1][int((w-1)/scale)][c]
19
20                 elif ((w-1)%scale==0):
21                     if (h%scale):
22                         image_zoom[h][w][c] = (h-1)/scale * image[0][int((w-1)/scale) - 1][c]
23                     else:
24                         image_zoom[h][w][c] = (1 - (h-1)/scale) * image[int((h-1)/scale) - 1][int((w-1)/scale) - 1][c] + (h-1)/scale * image[int((h-1)/scale)][int((w-1)/scale) - 1][c]
25
26                 else:
27                     if (h%scale and w%scale):
28                         image_zoom[h][w][c] = (w-1)/scale * (h-1)/scale * image[0][0][c]
29                     elif (h%scale):
30                         image_zoom[h][w][c] = (h-1)/scale * scale * ((1 - (w-1)/scale) * image[int((h-1)/scale) - 1][int((w-1)/scale) - 1][c] + (w-1)/scale * image[int((h-1)/scale) - 1][int((w-1)/scale)][c])
31                     elif (w%scale):
32                         image_zoom[h][w][c] = (w-1)/scale * scale * ((1 - (h-1)/scale) * image[int((h-1)/scale) - 1][int((w-1)/scale) - 1][c] + (h-1)/scale * image[int((h-1)/scale)][int((w-1)/scale) - 1][c])
33                     else:
34                         a = (1 - (h-1)/scale) * image[int((h-1)/scale) - 1][int((w-1)/scale) - 1][c] + (h-1)/scale * image[int((h-1)/scale)][int((w-1)/scale) - 1][c]
35                         b = (1 - (h-1)/scale) * image[int((h-1)/scale) - 1][int((w-1)/scale)][c] + (h-1)/scale * image[int((h-1)/scale)][int((w-1)/scale)][c]
36                         image_zoom[h][w][c] = (1 - (w-1)/scale) * a + (w-1)/scale * b
37
38     return image_zoom.astype(np.uint8)
```

```
def normalized_ssd(image1, image2):
    assert image1.shape == image2.shape, "Images must have the same shape"
    ssd = np.sum((image1 - image2) ** 2)

    num_pixels = image1.shape[0] * image1.shape[1]
    max_pixel_value = 255
    nssd = ssd / (num_pixels * max_pixel_value ** 2)

    return nssd
```

For image 1
 Normalized SSD using Nearest Neighbour: 0.0014433364007670365
 Normalized SSD using Bilinear Transform: 0.0014433364007670365

For image 2
 Normalized SSD using Nearest Neighbour: 0.0005491124941261907
 Normalized SSD using Bilinear Transform: 0.0005491124941261907



09.



```
#get image shape
h, w, _ = img_orig.shape

#define mask
mask = np.zeros((h, w), dtype=np.uint8)
print(h, w)
bgdModel = np.zeros((1, 65), np.float64)
fgdModel = np.zeros((1, 65), np.float64)

# rectangle to be given
rect = (50, 100, w, h-350)

cv.grabCut(img_orig, mask, rect, bgdModel, fgdModel, 7, cv.GC_INIT_WITH_RECT)

mask2 = np.where((mask == 2) | (mask == 0), 0, 1).astype('uint8')

img_foreground = img_orig * mask2[:, :, np.newaxis]

img_background = cv.subtract(img_orig, img_foreground)
```

```
kernel_size = 7
#blurred background = cv.GaussianBlur(img_background, (kernel_size, kernel_size), 0)
blurred_background = cv.blur(img_background, (kernel_size, kernel_size), 0)
```

Once the background image gets blur, edge near the missing flower gets closer values to 0 as most of the pixels are 0. Therefore, image gets little darker.

GitHub Link repo : https://github.com/Dulan24/S5_EN3061_Image-processing

Notebook link : https://github.com/Dulan24/S5_EN3061_Image-processing/blob/master/Activity%2001/answers/codes.ipynb